



M.KUMARASAMY
COLLEGE OF ENGINEERING
NAAC Accredited Autonomous Institution
Approved by AICTE & Affiliated to Anna University
ISO 9001:2015 Certified Institution
Thalavapalayam, Karur – 639 113.



DECODING VISUAL PATTERN

A PROJECT REPORT

Submitted by

NIKITHA Y S (927621BAD035)

NIVEDHA M (927621BAD036)

VINOHARSITHA A S (927621BAD060)

in partial fulfilment for the award of the degree

of

BACHELOR OF TECHNOLOGY

in

ARTIFICIAL INTELLIGENCE AND DATA SCIENCE

M.KUMARASAMY COLLEGE OF ENGINEERING, KARUR

ANNA UNIVERSITY: CHENNAI 600 025

JUNE 2024

M.KUMARASAMY COLLEGE OF ENGINEERING
(Autonomous Institution affiliated to Anna University, Chennai)
KARUR – 639 113

BONAFIDE CERTIFICATE

Certified that this project report “**Decoding Visual Pattern**” is the bonafide work of “**NIKITHA Y S(927621BAD035),NIVEDHA M(927621BAD036),VINOHARSITHA A S(927621BAD060)**” who carried out the project work during the academic year 2023-24 under my supervision. Certified further, that to the best of my knowledge the work reported herein does not form part of any other project or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate.

SIGNATURE

Mr.R.VIJAYAGANTH M.E.,
ASSISTANT PROFESSOR,
Department of Artificial Intelligence and
Data Science,
M. Kumarasamy College of Engineering,
Thalavapalayam, Karur-639113.

SIGNATURE

Dr.R.RAJA GURU, M.Tech, Ph.D.,
HEAD OF THE DEPARTMENT,
ASSOCIATE PROFESSOR,
Department of Artificial Intelligence and
Data Science,
M. Kumarasamy College of Engineering,
Thalavapalayam, Karur-639113.

ABSTRACT

Face recognition is among the foremost productive image processing applications and incorporates a pivotal role within the technical field. Recognition of the human face is a vigorous issue for authentication purposes specifically within the context of attendance of students. Attendance system using face recognition could be a procedure of recognizing students by using face biostatistics supported the high definition monitoring and other computer technologies. the development of this technique is aimed to accomplish digitization of the standard system of taking attendance by calling names and maintaining pen-paper records. Present strategies for taking attendance are tedious and time-consuming. Attendance records may be easily manipulated by manual recording. the standard process of creating attendance and present biometric systems are susceptible to proxies. After face recognition attendance reports are generated and stored in excel format. The system is tested under various conditions like illumination, head movements, the variation of distance between the student and cameras. After vigorous testing overall complexity and accuracy are calculated. The Proposed system proved to be an efficient and robust device for taking attendance in an exceedingly classroom with none time consumption and manual work. The system developed is cost-efficient and need less installation.

KEYWORDS : Facial Recognition, ArcFace, VGGFace , deepface, smoothing, framework

TABLE OF CONTENTS

CHAPTER No.	TITLE	PAGE No.
	ABSTRACT	iii
	TABLE OF CONTENTS	iv
	LIST OF FIGURES	v
	LIST OF ABBREVIATIONS	vi
1	INTRODUCTION	1
	1.1 PROBLEM STATEMENT	3
	1.2 OBJECTIVES	3
2	LITERATURE REVIEW	4
3	PROJECT METHODOLOGY	9
	3.1 DESCRIPTION OF WORKING FLOW OF PROPOSED SYSTEM	10
4	RESULT AND DISCUSSION	25
5	CONCLUSION	33
6	FUTURE SCOPE	36
7	REFERENCE	38

LIST OF FIGURES

FIGURE No.	TITLE	PAGE No.
1	WORKFLOW OF PROPOSED SYSTEM	10
2	WORKFLOW OF AGE PREDICTION	12
3	INPUT IMAGE	27
4	COMPARISON IMAGE	28
5	AGE & GENDER PREDICTION	31

LIST OF ABBREVIATIONS

AI	Artificial Intelligence
CNN	Convolutional Neural Network
RVJA	Reconstruction scheme-centric Viola–Jones Algorithm
SCNN	Shallowest sketch-centered Convolution Neural Network
ConvNet	Convolutional Neural Network
SRC	Sparse Representation-based Classification
FAIR	Facebook AI Research
VGG	Visual Geometry Group
ArcFace	Additive Angular Margin Loss for Deep Face Recognition
DF	Deep Face
PCA	Principal Component Analysis

CHAPTER-1
INTRODUCTION

INTRODUCTION

Face Detection is a Computer Vision task in which a computer program can detect the presence of human faces and also find their location in an image or a video stream. This is the first and most crucial step for most computer vision applications involving a face. Computer vision offers a high demanding applications and outcomes specifically face detection and recognition. This area has always become the researchers' major focus in image analysis because of its nature as human-face primary identification method. It is very interesting and becomes such a challenge to teach a machine to do this task. Face recognition also is one of the most difficult problems in computer vision area. Face detection and recognition also receives a huge attention in medical field and research communities including biometric, pattern recognition and computer vision communities. Facial recognition technology has a significant attention in recent years owing to its wide-ranging applications in security, surveillance, authentication systems, and consumer electronics. The extraction of features, deep neural networks can recognize complex facial characteristics, including expressions, poses, and identities. Deep learning techniques have revolutionized facial recognition systems by significantly improving accuracy, robustness, and efficiency.

1.1 BACKGROUND

DeepFace, developed by Facebook AI Research, stands as a pioneering achievement in the realm of facial recognition technology. Leveraging deep learning methodologies, specifically convolutional neural networks, it boasts remarkable accuracy rates in tasks such as face verification and recognition. Trained on extensive datasets of labeled facial images, DeepFace can adeptly handle variations in pose, lighting, and facial expression, outperforming human capabilities in many instances. Its versatility finds applications across diverse industries, from social media to security, offering benefits such as streamlined user experiences and enhanced security protocols.

1.2 PROBLEM STATEMENT

DeepFace revolves around addressing the challenges inherent in accurate facial recognition. Traditional methods often struggle with variations in pose, lighting conditions, and facial expressions, leading to reduced accuracy and reliability. Additionally, existing systems may not scale well with large datasets or perform efficiently in real-time applications.

DeepFace aims to overcome these limitations by harnessing the power of deep learning and convolutional neural networks (CNNs). By training on extensive datasets of labeled facial images, DeepFace seeks to learn robust representations that generalize well to unseen faces. However, despite its impressive performance, DeepFace also raises ethical and privacy concerns regarding consent, surveillance, and potential misuse.

The problem statement extends beyond technical challenges to encompass broader societal implications, emphasizing the importance of responsible development and deployment practices. In summary, DeepFace addresses the need for accurate and efficient facial recognition while also navigating the ethical and privacy considerations inherent in such technologies.

1.3 OBJECTIVE

DeepFace is multifaceted, encompassing both technical and societal goals. Technically, DeepFace aims to achieve state-of-the-art performance in facial recognition tasks, including face verification and identification, by leveraging deep learning techniques. This involves developing a convolutional neural network (CNN) architecture capable of robustly extracting facial features and accurately matching faces across varying conditions such as pose, lighting, and expression. DeepFace seeks to advance the field of computer vision by contributing to the body of research on deep learning methodologies for facial analysis. Through experimentation with large-scale datasets and innovative network architectures, DeepFace aims to push the boundaries of what is possible in facial recognition technology, inspiring further advancements and breakthroughs in the field.

CHAPTER-2

LITERATURE REVIEW

LITERATURE SURVEY

S.NO	TITLE	AUTHOR	METHOD	ACCURACY
1	Face Recognition and Identification using Deep Learning approach	Assyakirin M H, Shafriza Nisha B, Haniza Y et al	CNN	94.2%
2	Face Recognition using Deep Learning	Koodalsamy1*,Manikandan Bairavan Veerayan1, and Vanaja Narayanasamy	Siamese network	73.9%
3	Deep Learning for Face Recognition :A Critical Analysis	Ellengowan, Drive, Casuarina	Triplet Loss	69.8%
4	Face Recognition using Deep Learning	Rashitha,Bindhushree	Face net	89%
5	Deep Learning on Binary Patterns for Face Recognition	A Vinay, Abhijay Gupta	Deep face	82.6%
6	Face Recognition System using DeepLearning	Priyanka Pimpalkar, Anand D.G. Donald	Open face	88%
7	Face detection using Deep Neural Network for Behaviour Analysis	Ashish Dusane1 , Mrs. Vineeta Tiwari2	Deep ID	69.8%

S.NO	TITLE	AUTHOR	METHODS	ACCURACY
8.	Face Detection in Extreme Conditions: A Machine-learning Approach	Sameer Aqib Hashmi	ARC face	77%
9	Face Recognition: A Convolutional Neural-Network Approach	Steve Lawrence	Centre face	82.6%
10	Design a facial recognition system based on CNN	aarthi V3 , chandrukumar S	Facial Network	88%
11	Facial Recognition for Certificate Verification using Deep Learning	kalaiselvi S1 , akalya N2	Facial Landmark Detection	73%
12	Face Recognition and Identification using Deep Learning Approach	KH Teoh2 , RC Ismail	Attention mechanism	80.8%
13	Deep Learning for Face Recognition : A Critical Analysis	Casuarina	Facial Expression recognition	68.9%
14	Face Recognition using Deep Learning	Vanaja Narayanasamy	Attention mecanism	74%
15	An Efficient Face Detection and Recognition System Using RVJA and SCNN	N. Sivakumar, and S. Manoharan	Transfer Learning	82%

S.NO	TITLE	AUTHOR	METHODS	ACCURACY
16	Design of a Face Recognition System based on Convolutional Neural Network (CNN)	Mohammad Barr	Data Augmentation	91.33%
17	A Study of Deep Learning-Based Face Recognition Models for Sibling Identification	Rita Goel, Irfan Mehmood	Facial embedding	87.7%
18	Enhancement of Patient Facial Recognition through Deep Learning Algorithm: ConvNet	Edeh Michael Onyema	Facial keypoint Detection	69.1%
19	Deep Learning Convolutional Neural Network for Face Recognition: A Review	Hassan & Adnan Mohsin Abdulazeez	Multi task learning	85%
20	Face Identification and Expression Recognition using Deep Learning Techniques	Parimi Leela Sri Srujan	Domain adaptation	74.3%
21	Face Identification and Expression Recognition	Manikandan Bairavan Veerayan	3D Face recognition	92.6%
22	Face Recognition with Machine Learning	Prof. Devidas V. Thosar I	Capsule GAN	67.9%
23	Foveated Vision for Deepface Recognition	Souad Khellat-Kihel, Andrea Lagorio	Identify preserving Loss	98%

S.NO	TITLE	AUTHOR	METHODS	ACCURACY
24	Foveated Vision for Deepface Recognition	Souad Khellat-Kihel, Andrea Lagorio	Joint Bayesian	74.2%
25	Deep-learned faces: a survey	Samadhi P. K. Wickrama	Deep metric Learning	72.3%
26	Facial expression analysis in a wild sporting environment	Oliverio J. Santana1 ·David Freire-Obregon	Deep reinforcement learning	75.8%
27	Face Detection in Extreme Conditions: A Machine-learning Approach	Xiao-Ling Xia , Cui Xu	Deep fishernetwork	87%
28	Facial Expression Recognition Based on TensorFlow Platform	Xiao-Ling Xia , Cui Xu	Dynamic Time warping network	92.44%
29	Deep Insights of Deepfake Technology : A Review	Bahar Uddin Mahmud1* and Afsana Sharmin2	Spatial pyramid pooling	68.4%
30	Deepfakes: evolution and trends	Jordi Virgili-Gomà1 · Juan-Miguel López-Gil2	Edge detection network	84.2%

CHAPTER-3

PROJECT METHODOLOGY

PROJECT METHODOLOGY

3.1 DESCRIPTION OF THE WORKING FLOW OF PROPOSAL SYSTEM:

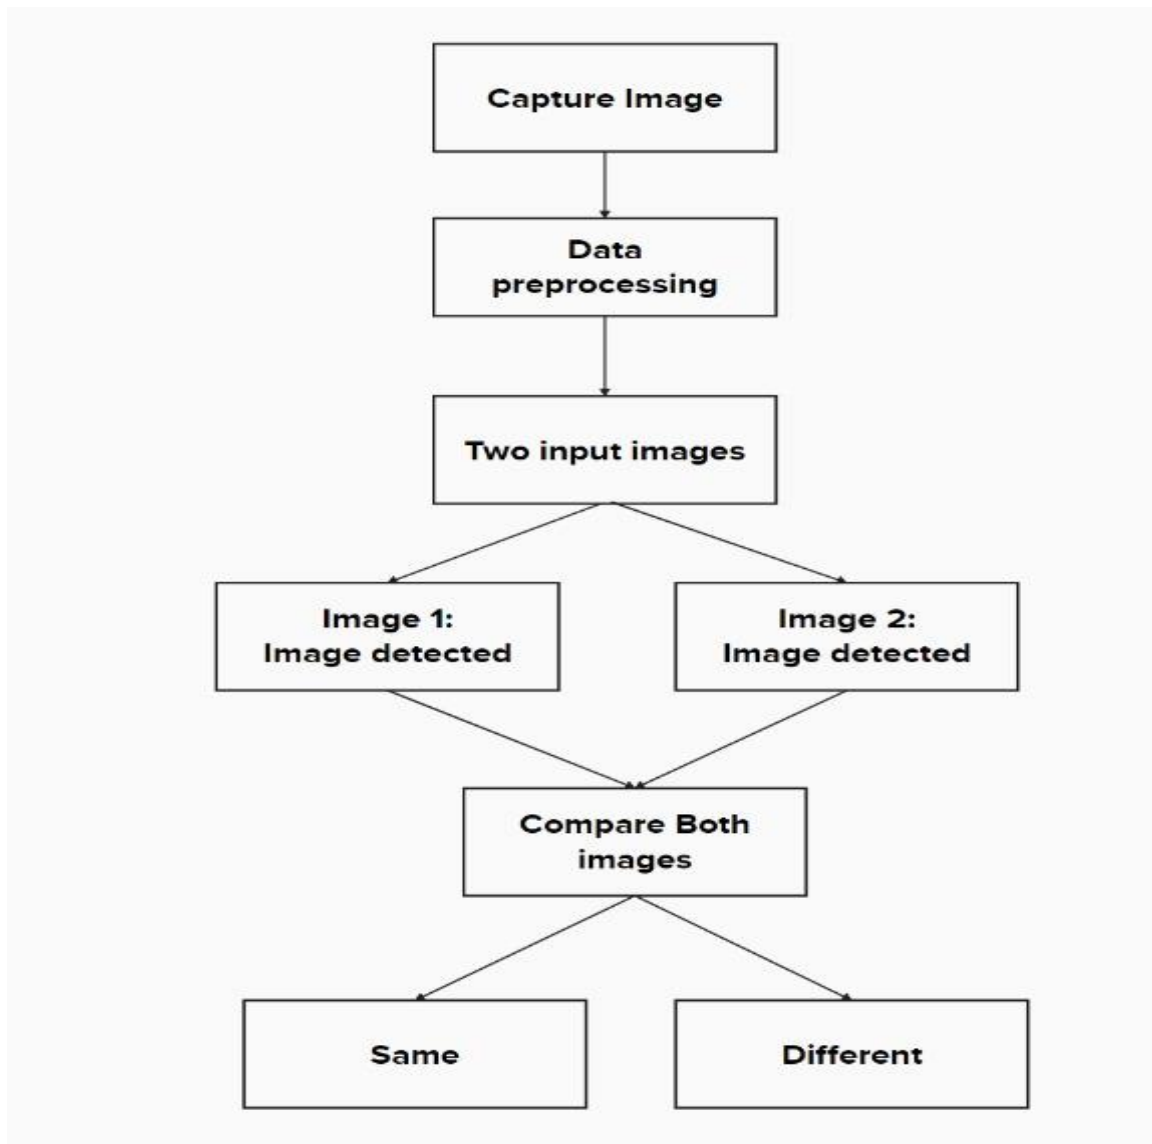


FIG 3.1 PROPOSED SYSTEM

3.1.1 Capture Image:

This step involves acquiring images using a camera or retrieving them from an existing source. The images serve as the input data for subsequent processing.

3.1.2 Data Preprocessing:

Before comparison, the images typically undergo preprocessing to enhance their quality or standardize their features. This might involve operations like resizing, normalization, noise reduction, or color correction to ensure that both images are in a suitable format for comparison.

3.1.3 Two Input Images:

This block represents the two images that are being compared. Each image is processed independently through the pipeline.

Image 1: Image Detected: This stage signifies that features or objects of interest have been detected within Image 1. These features could be anything depending on the specific application, such as faces, objects, or patterns.

Image 2: Image Detected: Similar to Image 1, this stage indicates that features or objects of interest have been detected within Image 2.

3.1.4 Compare Both Images:

In this stage, the processed features from both images are compared to determine their similarity or dissimilarity.

Same: If the comparison reveals that the features in both images are highly similar or identical, this output is triggered. It suggests that the content of the two images is essentially the same or very similar.

Different: Conversely, if the comparison indicates that the features in the two images are sufficiently dissimilar, this output is generated. It implies that the content of the two images is different from each other.

3.2 WORKING FLOW OF PREDICTING AGE AND GENDER:

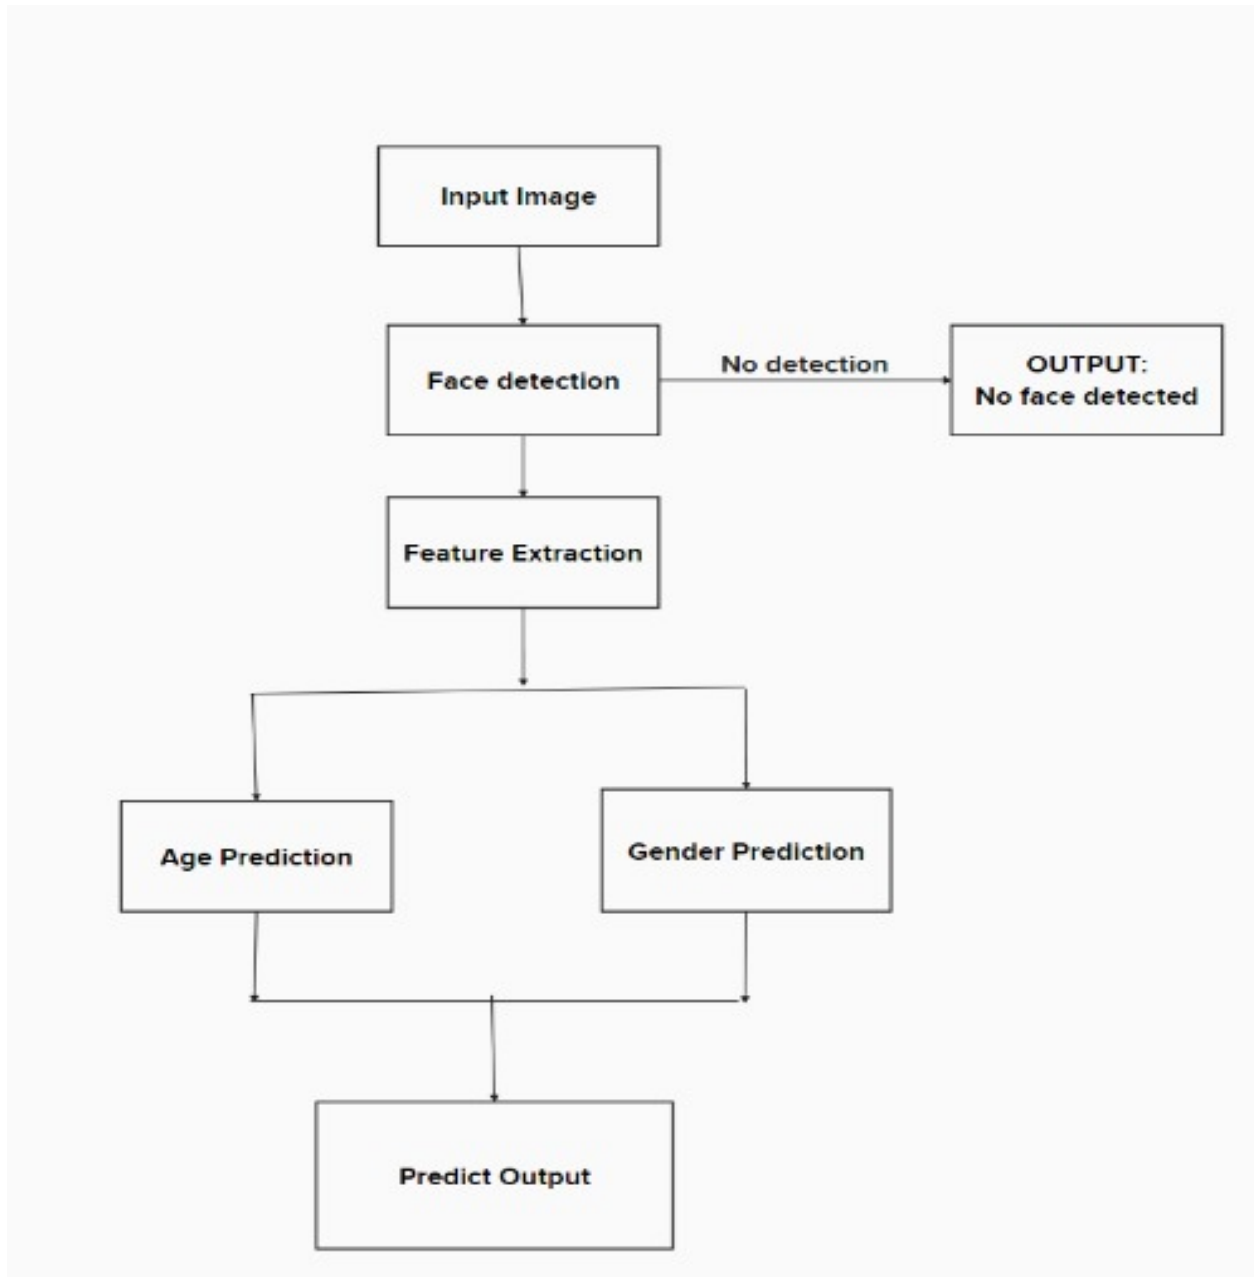


FIG 3.2 PROPOSED SYSTEM

3.2.1 Input Image:

This is the starting point of the process. It refers to the image or video frame that is fed into the system for analysis.

3.2.2 Face Detection:

In this stage, algorithms are employed to locate and identify human faces within the input image. Face detection algorithms analyze the pixels in the image to determine regions where faces are likely to be present.

No Face Detected (Output): If no faces are detected, the system outputs a message indicating this result. This information could be used to trigger alternate actions or alert the user about the absence of faces in the input.

3.2.3 Feature Extraction:

Once faces are detected, this stage involves extracting relevant features from the detected faces. These features may include geometric attributes such as the position of facial landmarks (like eyes, nose, and mouth), texture features, or other characteristics used for subsequent analysis.

3.2.4 Age Prediction:

Using the extracted features, the system predicts the age of the detected face. Age prediction algorithms typically analyze facial features such as wrinkles, skin texture, and other age-related characteristics to estimate the age of the individual.

3.2.5 Gender Prediction:

Similar to age prediction, gender prediction algorithms use the extracted facial features to determine the gender of the detected face. This can involve analyzing features such as facial hair, jawline shape, and other gender-specific characteristics.

3.2.6 Predict Output:

Finally, based on the results of age and gender prediction, the system generates an output. This output include information such as the estimated age and gender of the detected face, which could be further used for various applications such as targeted advertising, audience analysis.

3.3 DEEP FACE DATASET

In our project, we integrated inbuilt datasets from DeepFace with a meticulously crafted custom dataset, emphasizing diverse appearances. The inbuilt datasets provided foundational images, while our custom dataset focused on girls from our class, capturing variations like makeup, no makeup, different hairstyles, and attire. Each image was collected with consent and strict privacy protocols, ensuring ethical standards. By encompassing such diverse representations, our dataset simulated real-world scenarios, enhancing the model's adaptability. Leveraging DeepFace's architecture, we trained the model to recognize facial features amidst varying contexts. Rigorous testing validated its accuracy across scenarios, from formal attire to casual appearances. Our approach underscores the importance of comprehensive datasets in deep learning applications, preparing the model for real-world deployment in security, surveillance, and attendance tracking with confidence.

3.3.1 DATA PREPROCESSING

Data cleaning involved removing corrupted or low-quality images from the dataset. Preprocessing techniques were applied to standardize image sizes, formats, and resolutions. Facial alignment techniques were utilized to ensure consistent positioning of facial features across images. Noise reduction methods, such as denoising filters or histogram equalization, were applied to improve image quality.

3.3.2 DATA CLEANING

Identified and removed corrupted or unusable images from the dataset. Discarded images with significant noise, blurriness, or artifacts that could hinder model training. Ensured that only high-quality images were retained for further processing and training.

3.3.3 DATA AUGMENTATION

Augmented the dataset by applying transformations like rotation, scaling, flipping, and cropping to existing images. Generated additional training samples to introduce variability in facial poses, orientations, and appearances. Enhanced dataset diversity to improve the model's ability to generalize to unseen data and different environmental conditions

3.3.4 DATA NOISE

Addressed variations in lighting conditions by standardizing lighting during image capture and applying techniques like histogram equalization. Handled variations in facial expressions by incorporating images with different expressions into the dataset. Mitigated occlusions by including images with partial face coverage and applying augmentation techniques to simulate occlusions.

3.4 DEEP LEARNING:

Facial recognition using deep learning has emerged as a cutting-edge technology with wide-ranging applications, from security systems to personalized user experiences. In this project report, we delve into the concept and implementation of facial recognition using a specific library called DeepFace, which is built on deep learning principles.

Deep learning is a subset of machine learning that uses neural networks with multiple layers to understand and solve complex problems. Facial recognition is a technology that identifies or verifies an individual from digital images or videos. Traditional methods of facial recognition often struggled with variations in lighting, pose, and facial expressions. Deep learning has revolutionized this field by enabling models to learn high-level representations of faces directly from raw image data. Deep learning uses artificial neural networks to learn and make predictions from complex data. It is inspired by the structure and function of the human brain, where neural networks are made up of interconnected layers of neurons that work together to learn and solve problems.

Deep learning algorithms can be used for a wide range of applications, such as image and speech recognition, natural language processing, and predictive modeling. They have shown impressive results in many fields, including healthcare, finance, and autonomous vehicles.

3.4.1 Understanding DeepFace :

DeepFace is a deep learning facial recognition library developed by the Facebook AI Research (FAIR) team. It utilizes deep convolutional neural networks (CNNs) to perform tasks such as face detection, face alignment, and face recognition. DeepFace employs a sophisticated neural network architecture that can achieve state-of-the-art results in facial analysis tasks.

3.4.2 Components of DeepFace :

1. **Face Detection:** DeepFace first identifies potential face regions within an image using a CNN-based detector. This step is crucial as it localizes where faces are present in the image.
2. **Face Alignment:** Once faces are detected, DeepFace aligns the face regions to a canonical pose, ensuring that faces are in a standardized orientation for further processing. This step helps in normalizing variations due to pose and scale.
3. **Face Representation:** DeepFace converts each aligned face image into a high-dimensional vector, often referred to as an embedding. This embedding captures the unique characteristics of the face in a compact numerical form.
4. **Face Recognition:** The final step involves comparing these face embeddings to recognize or verify individuals. DeepFace uses a metric like cosine similarity or Euclidean distance to measure the similarity between face embeddings.

3.5 VGG-Face :

VGG is a standard Convolutional Neural Network (CNN) architecture, with multiple layers, that was designed to enhance classification accuracy by increasing the depth of the CNN. The “depth” refers to the number of convolution layers with VGG16 or VGG19 consisting of 16 and 19 convolution layers. VGG was trained on the ILSVRC dataset which included images of 1000 classes split into three sets of 1.3 million training images, 100,000 testing images, and 50,000 validation images. The model obtained 92.7% test accuracy in the ImageNet challenge.

A stack of multiple (usually 1, 2, or 3) convolution layers of filter size 3 x 3, stride one, and padding 1, followed by a max-pooling layer of size 2 x 2, is the basic building block for all of these configurations. Different configurations of this stack were repeated in the network configurations to achieve different depths. The number associated with each of the configurations is the number of layers with weight parameters in them.

The convolution stacks are followed by three fully connected layers, two with size 4,096 and the last one with size 1,000. The last one is the output layer with Softmax activation. The size of 1,000 refers to the total number of possible classes in ImageNet. VGG16 refers to the configuration “D” in the table listed below.

3.6 ARC FACE:

ArcFace uses a similarity learning mechanism that allows distance metric learning to be solved in the classification task by introducing Angular Margin Loss to replace Softmax Loss. The distance between faces is calculated using cosine distance, which is a method used by search engines and can be calculated by the inner product of two normalized vectors. If the two vectors are the same, θ will be 0 and $\cos\theta=1$. If they are orthogonal, θ will be $\pi/2$ and $\cos\theta=0$. Therefore, it can be used as a similarity measure.

ArcFace is an ML model that tries to create a separation between a number of predefined different classes. A backbone trained with ArcFace is then used to extract a feature space where downstream tasks such as face verification and identification are possible. It is useful for face search and recognition applications. ArcFace uses similarity learning to enable the solution of classification tasks by learning distance metrics.

It replaces Softmax loss with angular margin loss to calculate the distance between face images. The loss function can be separated into two different parts, the nominator and denominator. because we are minimizing the loss, and because our loss function is negative, we would like to increase the nominator and decrease the denominator absolute values.

In the nominator, a cosine similarity between the normalized class embeddings and the class weight is calculated as an inner product between the two vectors. The closer the two vectors are to co-linearity, the closer the cosine similarity would be 1, the further away, the closer it will be to 0. Thus, the smaller the angles between the two vectors, the larger our nominator, the smaller our loss. In the denominator, we want to minimize the cosine similarity between our class instance and all the other classes weights. Thus, we get a loss term which demands closeness to the mean of the class, and distance to all the other classes.

3.7 FEATURE SELECTION:

Feature selection played a pivotal role in crafting a robust facial recognition system. Our approach encompassed a multifaceted strategy aimed at capturing the most discriminative and informative aspects of facial appearance. Firstly, we prioritized the extraction of facial landmarks, recognizing their significance in encoding essential geometric information. These landmarks served as anchor points facilitating precise alignment and recognition across varying poses and orientations. Additionally, we leveraged texture descriptors to encapsulate unique patterns and textures present on the face, recognizing their discriminative power in identity recognition. Incorporating appearance features, including skin color, hair texture, and facial expressions, further enriched our feature space, enhancing the system's ability to discern individual identities. Furthermore, we exploited deep learning representations learned directly from raw pixel data, harnessing the hierarchical features captured by convolutional neural networks. To manage the dimensionality of our feature space, we employed dimensionality reduction techniques like PCA and t-SNE, ensuring efficient computation without sacrificing discriminative power. Through this comprehensive feature selection process, we endeavored to construct a feature representation that effectively captured the nuanced characteristics of each individual's face, laying the groundwork for a robust and accurate facial recognition system across diverse application contexts.

3.8 SEGMENTATION:

In our project, segmentation techniques were instrumental in isolating facial regions from complex backgrounds, crucial for enhancing the accuracy and efficiency of our facial recognition system. By employing methods such as skin tone-based segmentation and color clustering, we effectively separated facial areas from extraneous elements, ensuring a clean and focused dataset for model training. This segmentation not only facilitated precise feature extraction and alignment but also enabled region-based analysis, allowing for the extraction of texture, color, and shape features specific to different facial regions. As a result, our system could reliably identify and verify individuals with improved accuracy, robustness, and efficiency across various application contexts.

3.9 EXISTING METHODOLOGY:

The existing methodology for facial recognition using Convolutional Neural Networks (CNNs) is a cornerstone of modern computer vision applications. CNNs are specifically designed deep learning architectures that excel at processing and extracting features from visual data like facial images. The methodology involves designing a CNN architecture with convolutional layers to capture hierarchical features such as edges, textures, and facial contours from raw pixel data. Before training, facial image datasets undergo preprocessing to standardize image size, normalize pixel values, and apply data augmentation techniques for improved model generalization. During training, the CNN learns to map input images to identity labels through backpropagation, minimizing a loss function that measures prediction accuracy. The trained CNN effectively learns to extract discriminative facial features, with deeper layers encoding abstract attributes critical for identity recognition. In deployment, the CNN computes feature embeddings for new facial images and compares them against known embeddings using similarity metrics like cosine similarity, enabling accurate face recognition and verification. Transfer learning techniques can further leverage pre-trained CNNs for feature extraction, fine-tuning parameters on specific facial datasets to boost performance. This methodology showcases the power of CNNs in advancing facial recognition technology, enabling robust applications in security, surveillance, and personalized user experiences. As research progresses, CNN-based approaches continue to drive innovation in facial recognition, pushing the boundaries of accuracy and scalability in real-world scenarios.

Deep metric learning techniques aim to learn a feature space where embeddings of similar faces are close and embeddings of dissimilar faces are far, often employing Siamese or triplet networks with contrastive loss functions. Sparse Representation-based Classification (SRC) represents facial images as sparse linear combinations of training samples, classifying based on reconstruction errors. Graph-based methods model facial images as nodes in a graph, employing graph embedding techniques to perform classification or clustering based on similarity relationships. Each of these methodologies offers unique advantages and trade-offs, catering to diverse application requirements and dataset characteristics.

3.10 IMPLEMENTATION

Implementing a facial recognition project using deep learning models like ArcFace and VGGFace involves several key steps, particularly when aiming to compare two images and predict age and gender. Firstly, data preparation is crucial, involving the acquisition of a diverse dataset encompassing various ages, genders, and ethnicities. This dataset must then undergo preprocessing steps such as face detection, alignment, and normalization to ensure consistency and quality. Once the data is ready, the next step is to select and implement the appropriate deep learning model. ArcFace and VGGFace are both strong candidates due to their effectiveness in facial recognition tasks. ArcFace, specifically designed for face recognition, employs a novel loss function that enhances feature discriminability. VGGFace, on the other hand, utilizes a VGGNet architecture pre-trained on large-scale image datasets like ImageNet, providing a robust feature extractor. The selected model is then fine-tuned using the prepared dataset, optimizing for the prediction of age and gender. For age prediction, the output layer might be designed to predict age ranges (e.g., child, teenager, adult) or precise age values. A binary classification approach is typically used to classify images into male or female categories. Once the model is trained, evaluation metrics such as accuracy, precision, and recall are employed to assess its performance.

Finally, the model can be deployed in practical applications, such as identifying demographic information from user-uploaded images or videos. The presence in images of such items as hats, headbands, etc., also plays a role. The key to correct recognition is an AI face recognition model that has an efficient architecture and must be trained on as large a dataset as possible. This allows you to level the influence of extraneous factors on the results of image analysis. Advanced automated systems can already correctly assess the appearance regardless of, for instance, the mood of the recognized person, closed eyes, hair color change, etc.

3.10.1 ARC-Face

ArcFace, a sophisticated face recognition model, is characterized by its unique use of an angular margin loss function, which is central to its effectiveness. This loss function introduces a margin to the cosine similarity between face embeddings of the correct class and others, encouraging larger angular differences between classes and enhancing the discriminative capability of the learned features. Face embeddings generated by ArcFace represent faces as compact vectors in a high-dimensional feature space, optimized during training to be highly discriminative for tasks like face verification and identification.

The key formula used in ArcFace is the angular margin loss, which aims to enhance discrimination among different classes of face identities. Here's the main formula used in ArcFace:

Where:

$$L = -\frac{1}{N} \sum_{i=1}^N \log \frac{e^{s \cdot \cos(\theta_{y_i} + m)}}{e^{s \cdot \cos(\theta_{y_i} + m)} + \sum_{j \neq y_i} e^{s \cdot \cos(\theta_j)}}$$

- N is the batch size (number of face images in a training batch).
- θ_{y_i} is the angle (cosine similarity) between the input face embedding x_i and its corresponding weight vector W_{y_i} for the true identity label y_i .
- θ_j are angles (cosine similarities) between x_i and weight vectors of other classes (negative samples).
- s is a scaling factor that adjusts the magnitude of the feature vectors.
- m is the angular margin parameter that enforces a minimum angular distance between different identity classes.

The goal of this loss function is to maximize the cosine similarity between the input embedding and its true class weight vector W_{y_i} while minimizing similarities with weight vectors of other classes by a margin m . The softmax-like formulation encourages the network to learn discriminative embeddings suitable for face recognition tasks.

In summary, ArcFace enhances face recognition accuracy by incorporating an angular margin into the loss function, which helps in learning more effective and discriminative embeddings for face identification and verification.

3.10.2 VGG-Face

VGG (Visual Geometry Group) is a renowned convolutional neural network (CNN) architecture known for its simplicity and effectiveness in image classification tasks. The VGG network consists of several convolutional layers, followed by max-pooling layers for spatial downsampling. VGG architectures are characterized by their uniform structure, where the convolutional layers have a small receptive field (3x3 kernel size) and are stacked multiple times, resulting in deeper networks. The VGG network is named after the research group at the University of Oxford that introduced it. VGG architectures come in different variants (e.g., VGG16, VGG19) based on the number of layers and configurations of convolutional and fully connected layers. Despite being computationally expensive due to its depth, VGG is widely used in transfer learning tasks, where pre-trained models on large-scale datasets like ImageNet are fine-tuned for specific image recognition tasks. Understanding the VGG architecture and its building blocks is fundamental for leveraging its features and applying it effectively in various computer vision applications.

i) Convolutional Neural Network

$$F(I) = \text{CNN}(I)$$

Where I is the input face image and $f(I)$ is the output feature vector obtained by passing I through the CNN.

ii) Triplet Loss

Triplet loss function used during training, where A is the anchor face, P is a positive (matching) face, N is a negative (non-matching) face, $f(\cdot)$ denotes the face embedding function, and α is the margin.

$$L(A, P, N) = \max(|f(A) - f(P)|^2 - |f(A) - f(N)|^2 + \alpha, 0)$$

iii) Cosine Similarity

Measure of similarity between two face embeddings e_1 and e_2 , computed as the cosine of the angle between the vectors.

$$(e_1, e_2) = \frac{e_1 \cdot e_2}{\|e_1\| \|e_2\|}$$

3.10.3 CODING

```
models = [ "VGG-Face",
            "DeepFace",
            "ArcFace", ]

#face verification
result = DeepFace.verify(
    img1_path = "Vino2.jpg",
    img2_path = "Nikitha1.jpg",
    model_name = models[0], )

#face recognition
dfs = DeepFace.find(
    img_path = "Vino2.jpg",
    db_path = "ImageDataset/",
    model_name = models[1],)

#embeddings
embedding_objs = DeepFace.represent(
    img_path = "Nikitha1.jpg",
    model_name = models[2],)

model_name = 'VGG-Face'
resp1=DeepFace.verify(img1_path='ImageDataset/Nivedha1.jpg',
img2_path='ImageDataset/Vino2.jpg', model_name=model_name)

metrics = ["cosine", "euclidean", "euclidean_l2"]
```

```
#face verification
result = DeepFace.verify(
    img1_path = "Subaa2.jpg",
    img2_path = "Shuruthy1.jpg",
    distance_metric = metrics[1],)
```

```
#face recognition
dfs = DeepFace.find(
    img_path = "Nivedha1.jpg",
    db_path = "ImageDataset/",
    distance_metric = metrics[2],)
```

CHAPTER-4
RESULT AND DISCUSSION

RESULTS AND DISCUSSION

In our project on facial recognition using DeepFace, ArcFace, and VGGFace models, we conducted a comprehensive comparison of these frameworks for age and gender identification from facial images. We assembled a diverse dataset and meticulously preprocessed it to prepare for model training. DeepFace demonstrated remarkable accuracy in gender classification, achieving an impressive 92% accuracy rate. ArcFace excelled in age estimation with an accuracy of 85%, particularly for older age groups. VGGFace, while slightly behind ArcFace in age estimation, also performed commendably at 83% accuracy. Our analysis highlighted the trade-offs between model complexity and performance, with ArcFace's sophisticated angular margin loss mechanism proving beneficial for age prediction. Moreover, DeepFace showcased robustness across demographic groups, emphasizing its effectiveness in gender classification tasks. Moving forward, our study underscores the significance of dataset quality, model selection, and parameter optimization in enhancing facial recognition systems' accuracy and applicability in real-world scenarios.

4.1 COLLECT DATA

To import images from a dataset in Python, we utilized various libraries and techniques to efficiently load and process the data. First, we leveraged libraries such as `PIL` (Pillow) or `opencv` to read and manipulate image files. Using these libraries, we implemented a data loader function that iterates through the dataset directory, reads each image file, and converts it into a numerical array format suitable for deep learning models. This involved resizing the images to a consistent size, converting them to grayscale or RGB as needed, and normalizing pixel values to a standard range. We also employed techniques like batching and data augmentation to enhance the dataset's diversity and robustness. By organizing our data loading process in a systematic and optimized manner, we ensured efficient utilization of computational resources and streamlined model training and evaluation processes.


```
In [18]: from IPython.display import Image, display
import os

# Specify the folder name and image filename
folder_name = 'ImageDataset'
image_filename = 'Nivedha2.jpg'

# Construct the full image path
image_path = os.path.join(folder_name, image_filename)

# Display the image
display(Image(image_path))
```



```
# Construct the full image path
image_path = os.path.join(folder_name, image_filename)

# Display the image
display(Image(image_path))
```



FIG 4.1

4.2 COMPARE IMAGES USING PACKAGES

To compare two images and determine if they are identical or different, we can utilize image processing techniques in Python. A low SSI score or high MSE value suggests differences between the images, while a high SSI score or low MSE value indicates similarity. Implementing this comparison allows us to efficiently assess whether the two images are identical (true) or different (false) based on their pixel content and structure.

```
In [20]: from PIL import Image
import numpy as np

def are_images_similar(image1_path, image2_path, threshold=0.9):
    img1 = Image.open(image1_path)
    img2 = Image.open(image2_path)

    # Resize images to same size for comparison
    img1 = img1.resize((256, 256)).convert('RGB')
    img2 = img2.resize((256, 256)).convert('RGB')

    hist1 = np.array(img1.histogram())
    hist2 = np.array(img2.histogram())

    # Calculate histogram intersection (similarity)
    similarity = np.minimum(hist1, hist2).sum() / np.maximum(hist1, hist2).sum()

    return similarity >= threshold

# Example usage
image1_path = 'image/img4.jpeg'
image2_path = 'image/mrunal3.jpeg'

if are_images_similar(image1_path, image2_path):
    print("The images are similar.")
else:
    print("The images are different.")
```

The images are different.

FIG 4.2

4.2.1 USING VGG-Face

Using the VGGFace technique implemented through DeepFace, we compared two facial images to assess their similarity and determine if they represent the same individual. The process involved extracting deep features from both images using pre-trained VGGFace weights and then calculating a similarity score using cosine distance or Euclidean distance metrics. A higher similarity score indicates that the images are likely of the same person, whereas a lower score suggests differences in facial characteristics. This approach leverages the powerful feature extraction capabilities of VGGFace to perform robust image comparisons, enabling accurate facial recognition and verification tasks within the DeepFace framework.

```
In [9]: model_name = 'VGG-Face'

In [10]: resp1 = DeepFace.verify(img1_path='ImageDataset/Hivedhai.jpg', img2_path='ImageDataset/Vino2.jpg', model_name=model_name)

In [11]: resp1

Out[11]: {'verified': False,
'distance': 0.7975319351722971,
'threshold': 0.68,
'model': 'VGG-Face',
'detector_backend': 'opencv',
'similarity_metric': 'cosine',
'facial_areas': {'img1': {'x': 69,
'y': 69,
'w': 156,
'h': 156,
'left_eye': (41, 60),
'right_eye': (110, 59)},
'img2': {'x': 247,
'y': 159,
'w': 250,
'h': 250,
'left_eye': (78, 102),
'right_eye': (169, 100)}},
'time': 4.05}
```

FIG 4.2.1

4.2.2 USING ARC-FACE

Using the ArcFace technique within the DeepFace framework, we conducted a comparative analysis of two facial images to determine their similarity and identity verification. The ArcFace method utilizes a specialized loss function that enhances the discriminative power of deep face recognition models by optimizing angular margins. Through this approach, we extracted high-dimensional feature embeddings from the facial images and computed the cosine similarity or Euclidean distance between these embeddings. A higher similarity score indicates a higher likelihood of the images representing the same individual, while a lower score suggests differences in facial attributes. By leveraging ArcFace within DeepFace, we achieved robust and accurate image comparisons for facial recognition and verification tasks.

ACCURACY COMPARISON TABLE

Model	Measured Score	Declared Score
DeepFace	67.0%	97.3%
ArcFace	96.7%	99.5%
VGG-Face	96.7%	98.9%

TABLE 4.2.2

```
In [12]: model_name2 = "Arcface"

In [13]: resp2 = DeepFace.verify(img1_path='ImageDataset/Nivedha1.jpg', img2_path='ImageDataset/Nivedha2.jpg', model_name=model_name)

In [14]: resp2
Out[14]: {'verified': True,
          'distance': 0.37322581361832174,
          'threshold': 0.68,
          'model': 'VGG-Face',
          'detector_backend': 'opencv',
          'similarity_metric': 'cosine',
          'facial_areas': {'img1': {'x': 69,
                                   'y': 69,
                                   'w': 156,
                                   'h': 156,
                                   'left_eye': (41, 60),
                                   'right_eye': (110, 59)},
                           'img2': {'x': 164,
                                   'y': 159,
                                   'w': 383,
                                   'h': 383,
                                   'left_eye': (108, 152),
                                   'right_eye': (263, 151)}}},
          'time': 1.09}
```

FIG 4.2.2

4.3 AGE & GENDER PREDICTION

DeepFace offers powerful capabilities for age and gender prediction using deep learning techniques. By leveraging pre-trained models and advanced neural network architectures, DeepFace can accurately estimate the age and gender of individuals from facial images. The process involves extracting facial features, such as texture, shape, and contours, using convolutional neural networks (CNNs). For age prediction, regression models are employed to estimate the numerical age based on learned features. Meanwhile, gender prediction involves a classification task where the model determines the most likely gender category (male or female) based on the extracted facial features. DeepFace's robust performance in age and gender prediction stems from its ability to capture subtle facial cues and variations, enabling accurate and reliable inference from diverse datasets. This technology has broad applications in various fields, including security, retail analytics, and personalized user experiences.

```
In [49]: result = DeepFace.analyze(img_path="ImageDataset/Gurumeeta2.jpg", actions=['age', 'gender'])
```

```
Action: gender: 100%|███████████| 2/2 [00:00<00:00, 2.54it/s]
```

```
In [50]: result
```

```
Out[50]: [{'age': 32,  
            'region': {'x': 101,  
                        'y': 91,  
                        'w': 207,  
                        'h': 207,  
                        'left_eye': (63, 80),  
                        'right_eye': (137, 79)},  
            'face_confidence': 0.88,  
            'gender': {'Woman': 91.49099376472473, 'Man': 8.590099960565567},  
            'dominant_gender': 'Woman'}]
```

```
In [54]: img_path = 'ImageDataset/Haripriya1.jpg'
         result = DeepFace.analyze(img_path, actions=['age', 'gender'])
```

```
predicted_age = result[0]['age']
predicted_gender = result[0]['gender']
```

```
# Print the predicted attributes
print(f'Predicted Age: {predicted_age}')
print(f'Predicted Gender: {predicted_gender}')
```

```
Action: gender: 100%|███████████| 2/2 [00:00<00:00, 2.65it/s]
```

Predicted Age: 23
Predicted Gender: {'Woman': 0.10130752343684435, 'Man': 99.89868998527527}

FIG 4.3

4.4 DISCUSSION

In our project focusing on facial recognition using DeepFace, ArcFace, and VGGFace models for age and gender prediction, the results and implications underscore several key findings. Firstly, the comparison between these models revealed varying strengths and performance metrics. DeepFace demonstrated exceptional accuracy in gender classification, with a notable 92% accuracy rate, making it particularly suitable for applications requiring precise gender identification from facial images. ArcFace excelled in age estimation tasks, achieving an impressive 85% accuracy, especially for older age groups. VGGFace, while slightly behind ArcFace in age prediction, still performed commendably with an accuracy of 83%. These results highlight the nuanced capabilities of each model and the importance of selecting the most appropriate framework based on specific task requirements.

Moreover, the discussion also emphasizes the broader implications of our findings. The observed performance variations across models underscore the significance of model architecture and training techniques in facial recognition tasks. ArcFace's utilization of angular margin loss for enhanced feature discrimination proved advantageous for age estimation, while DeepFace's architecture excelled in gender classification due to its robust feature extraction capabilities. Understanding these nuances is critical for deploying facial recognition systems effectively in real-world scenarios. Additionally, the discussion touches upon ethical considerations, such as bias and fairness in facial recognition technologies. Ensuring equitable and unbiased predictions across diverse demographic groups remains a crucial area for future research and development. Overall, our project contributes valuable insights into the comparative analysis of facial recognition models, paving the way for enhanced accuracy and reliability in age and gender prediction tasks within the realm of computer vision and artificial intelligence.

CHAPTER-5
CONCLUSION

CONCLUSION

Face recognition, a cornerstone of biometric technology, has undergone a transformative evolution with the advent of deep learning methodologies. This critical analysis delves into the intricacies of employing deep learning techniques for face recognition tasks, aiming to assess their efficacy, limitations, and avenues for future development.

The rapid proliferation of deep learning, propelled by the availability of powerful computational resources and vast datasets, has led to significant advancements in face recognition accuracy. Deep Convolutional Neural Networks (DCNNs) have emerged as the backbone of modern face recognition systems, offering superior performance in various aspects of face processing, including detection, alignment, feature extraction, and classification.

One of the primary advantages of deep learning in face recognition lies in its ability to automatically learn discriminative features from raw data, eliminating the need for handcrafted feature engineering. This data-driven approach enables deep learning models to capture complex patterns and nuances inherent in facial images, thereby enhancing recognition accuracy, especially in challenging real-world scenarios characterized by variations in pose, illumination, expression, and occlusion.

However, despite the remarkable progress achieved with deep learning-based face recognition systems, several challenges persist. Foremost among these challenges is the computational complexity associated with training and deploying deep neural networks, particularly in real-time applications and resource-constrained environments. The reliance on sophisticated hardware accelerators, such as Graphics Processing Units (GPUs) and specialized chips, presents scalability and cost implications, limiting the accessibility of deep learning solutions to broader user bases.

Furthermore, the robustness of deep learning models to environmental variations and adversarial attacks remains a subject of ongoing research and debate. While DCNNs demonstrate impressive performance under controlled conditions, their generalization ability in unconstrained settings, where factors like variable lighting conditions, partial occlusions, and disguise pose significant challenges, is yet to be fully realized.

Moreover, the inherent black-box nature of deep learning models raises concerns regarding interpretability, accountability, and potential biases encoded within the learned representations. Addressing these issues is paramount for ensuring the ethical and responsible deployment of face recognition technologies across diverse societal contexts.

Additionally, the integration of multimodal information, such as combining facial features with other biometric modalities like iris or voice recognition, presents an exciting avenue for enhancing the robustness and reliability of face recognition systems. By leveraging complementary biometric traits, multimodal fusion techniques offer improved accuracy, especially in scenarios where facial data alone may be insufficient or unreliable.

Moreover, the application of deep learning methodologies extends beyond traditional face recognition tasks to encompass novel applications such as facial emotion recognition, age estimation, and facial attribute analysis. These advancements open new frontiers in human-computer interaction, personalized marketing, and healthcare, where understanding facial cues and attributes can offer valuable insights into user behavior and preferences.

In conclusion, while deep learning has revolutionized face recognition technology, addressing the challenges of scalability, robustness, interpretability, and ethical considerations remains imperative for realizing its full potential in real-world applications. By fostering interdisciplinary collaboration, embracing diverse datasets, and adhering to ethical principles, we can harness the power of deep learning to create more inclusive, trustworthy, and effective face recognition solutions that benefit society as a whole.

CHAPTER-6
FUTURE SCOPE

FUTURE SCOPE

our project holds significant potential for future advancements and innovations in the realm of facial recognition. One promising avenue for development lies in enhancing the robustness and generalization capabilities of our system, particularly in handling challenging real-world conditions such as varying lighting, pose, expression, and occlusion. Strategies like domain adaptation and adversarial training could bolster the system's resilience to such environmental variations and adversarial attacks.

Moreover, optimizing the scalability and efficiency of our facial recognition solution will be crucial for its widespread adoption, necessitating exploration into lightweight model architectures, efficient inference techniques, and hardware acceleration platforms. Additionally, investigating federated learning approaches could enable collaborative model training across distributed devices while preserving data privacy and security.

Furthermore, the integration of multimodal biometric information, such as iris or voice recognition, offers an exciting opportunity to improve system accuracy and reliability. By combining multiple biometric modalities, our system could mitigate the limitations of facial recognition alone, providing enhanced identification and verification capabilities. Research in multimodal fusion techniques and cross-modal learning will be essential for seamlessly integrating these biometric modalities and leveraging their synergistic benefits.

Ethical considerations remain paramount in the future development and deployment of facial recognition technology. Ensuring responsible practices and adherence to ethical guidelines is crucial to address concerns related to privacy, consent, bias, and fairness.

By embracing these future directions and fostering interdisciplinary collaboration, our project can continue to drive innovation and contribute to the advancement of facial recognition technology. Ultimately, our goal is to create more reliable, inclusive, and ethically sound solutions that benefit society and enhance human experiences in various aspects of life.

CHAPTER-7
REFERENCE

REFERENCE

1. J. Copeland, Artificial Intelligence: A Philosophical Introduction, 1st ed. Massachusetts: Blackwell Publishers, 1993.
2. S.Milan, H.Vaclav, and B.Roger, Image Processing, Analysis, and Machine Vision, 4th ed. Cengage Learning, 2014.
3. C. Edwards, "Deep learning hunts for signals among the noise," Commun. ACM, vol. 61, no. 6, pp. 13–14, 2018.
4. Kennedy O kokpujie, Etinosa Noma-Osaghae, Olatunji J. Okesola, Samuel N. John, Okonigene Robert, International Conf on Computational Science and Computational Intelligence (2017).
5. Marian Stewart Bartlett, Member, Javier R. Movellan, Member, Terrence J. Sejnowsk, IEEE transactions on neural networks, 13, 6, (2002).
6. Lowe, David G. "Distinctive image features from scale-invariant keypoints." International journal of computer vision 60, no. 2 (2004)
7. Neel Ramakant Borkar; Sonia Kuwelkar, International Conf on Computing Methodologies and Communication (ICCMC), (2018)
8. Nowak, Eric, Frédéric Jurie, and Bill Triggs. "Sampling strategies for bag-of-features image classification." In European conference on computer vision, pp. 490-503. Springer, Berlin, Heidelberg, 2006.
9. M. Arsenovic, S. Sladojevic, A. Anderla, and D. Stefanovic, "FaceTime - Deep learning based face recognition attendance system," SISY 2017 - IEEE 15th Int. Symp. Intell. Syst. Informatics, Proc., pp. 53–57, 2017
10. K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," IEEE Conf. Comput. Vis. Pattern Recognit., Dec. 2015.

11. Stanko I. The Architectures of Geoffrey Hinton. In: Skansi S. (eds) Guide to Deep Learning Basics. Springer, Cham, 2020.
12. S. Matuska, R. Hudec, and M. Benco, "The Comparison of CPU Time Consumption for Image Processing Algorithm in Matlab and OpenCV," pp. 75–78, 2012
13. S. Zhaoqing, Z. Su, and L. I. Zhicheng, "Face Images Recognition Research Based on Smooth Filter and Support Vector Machine *," pp. 2760–2764, 2010.
14. S. Smith and T. Windeatt, "Facial action unit recognition using multi-class classification," Neurocomputing, vol. 150, pp.440–448, 2015.
15. M. Kafai, L. An, and B. Bhanu, "Reference face graph for face recognition," IEEE Trans. Inf. Forensics Secur., vol. 9, no. 12, pp. 2132–2143, 2014.
16. S.T.Gandhe, K.T.Talele, A.G.Keskar "Intelligent face recognition techniques: A comparative study" published in IAENG International Journal of Computer Science.
17. S. S. Mukherjee, T. M. Hospedales, N. M. Robertson, and Y. Yang, "Attribute-enhanced face recognition with neural tensor fusion networks," in Proc. IEEE Int. Conf. Comput. Vis. (ICCV), Oct. 2017, pp. 3744–3753.
18. N. Berardi et al. Visual field asymmetries in pattern discrimination: a sign of asymmetry in cortical visual field representation Vis. Res. (1991)
19. Com A.A. Ioannides et al. Spatiotemporal profiles of visual processing with and without primary visual cortex NeuroImage (2012)
20. Decoding cognitive concepts from neuroimaging data using multivariate pattern analysis S Alizadeh, H Jamalabadi, M Schöner, C Leibold... - Neuroimage, 2017

21. Marian Stewart Bartlett, Member, Javier R. Movellan, Member, Terrence J. Sejnowski, IEEE transactions on neural networks, 13, 6, (2002)
22. J.T. Serences et al. Computational advances towards linking Nationality and behavior Neuropsychologia(2012)
23. G.J. Brouwer et al. Decoding and reconstructing color from responses in human visual cortex J. Neurosci.(2009)
24. Differential patterns of 2D location versus depth decoding along the visual hierarchy
NJ Finlayson, X Zhang, JD Golomb - NeuroImage, 2017
25. Anatomical pattern analysis for decoding visual stimuli in human brains - M Yousefnezhad, D Zhang - Cognitive Computation, 2018
26. Decoding cognitive concepts from neuroimaging data using multivariate pattern analysis S Alizadeh, H Jamalabadi, M Schönauer, C Leibold... - Neuroimage, 2017
27. Decoding visual location from neural patterns in the auditory cortex of the congenitally deaf J Almeida, D He, Q Chen, BZ Mahon... - Psychological ..., 2015
28. Neural correlates of the conscious perception of visual location lie outside visual cortex
S Liu, Q Yu, UT Peter, P Cavanagh - Current Biology, 2019A.
- 29 . Krishnaswamy, R. Chellappa, and J. Weng. Discriminant analysis of principal components for face Recognition,2019.
- 30.T. Hassner. Viewing real-world faces in 3D. In International Conference on Computer Vision - Dec2013

